

Assessing pre-training bias in Health data and estimating its impact on machine learning algorithms

Diego Dimer Rodrigues¹, Mariana R Mendoza¹

¹Instituto de Informática – Universidade Federal do Rio Grande do Sul (UFRGS)
Caixa Postal 15.064 – 91.501-970 – Porto Alegre – RS – Brazil

{ddrodrigues, mrmendoza}@inf.ufrgs.br

Abstract.

1. Introduction

Machine learning (ML) is increasingly used in various domains, including Health, where it assists in diagnosis, prognosis, and treatment [Chen et al. 2021]. However, ML algorithms can potentially introduce biases, leading to unfair and potentially harmful outcomes [Obermeyer et al. 2019, Juhn et al. 2022].

This work aims to measure the risk of bias in Health datasets using pre-training bias metrics and assesses their impact on the performance of supervised ML algorithms, correlating the bias metrics to the overall models' performance. We analyze four pre-training bias metrics [Hardt et al. 2021]: Class Imbalance (CI), Kullback-Leibler (KL) Divergence, Kolmogorov-Smirnov (KS), and Conditional Demographic Disparity in Labels (CDDL).

1.1. Class Imbalance (CI)

CI measures the representation of different classes within a dataset. It ranges from -1 to 1, with values near zero indicating balanced representation. Positive values indicate overrepresentation of the advantaged group, while negative values indicate overrepresentation of the disadvantaged group.

1.2. Kullback-Leibler (KL) Divergence

KL Divergence measures the divergence between the label distributions of two facets. It ranges from 0 to $+\infty$, with values near zero indicating similar distributions and higher values indicating greater divergence.

1.3. Kolmogorov-Smirnov (KS)

KS measures the maximum divergence between labels in the distribution for different facets. It ranges from 0 to 1, with values near zero indicating evenly distributed labels and values near one indicating unbalanced labels.

1.4. Conditional Demographic Disparity in Labels (CDDL)

CDDL measures the proportion of negative outcomes in a specific facet of a dataset. It ranges from -1 to 1, with positive values indicating demographic disparity favoring the advantaged group and negative values indicating disparity favoring the disadvantaged group.

2. Related Work

Many libraries and open-source tools have been developed to address machine learning fairness. As an example, AIF 360 library [Bellamy et al. 2018] provides different techniques to identify and mitigate bias, Studies have shown that word embedding algorithms may encode marginalized populations differently, perpetuating social bias. For example, [Zhang et al. 2020] demonstrated that such models associated African Americans and Blacks with prison and Whites and Caucasians with hospital in a fill-in-the-blank style course of action. Similarly, Fairness Measures [Zehlike et al. 2017] provides different ways of classifying fairness, bias and discrimination for any *csv* dataset. Additionally, FairLens [Synthesized.io 2023] is a python library that automatically discover and measures bias in data, providing a report on the data quality.

AI models trained on health data can be biased due to protected attributes, such as ethnicity and race. Various studies have explored different approaches for evaluating and mitigating bias in these models. [Júnior et al. 2022] analyzed the COMPAS dataset on criminal recidivism, and found that the classifier was more prone to predict recidivism for Black people, favoring other demographics. In [Noseworthy et al. 2020], the author recommended that AI tools report the model’s performance for all the protected attributes. The study in [Park et al. 2021] found that bias mitigation algorithms outperformed simply removing the protected attribute that might cause bias. [Mandhala et al. 2022] used pre-training bias metrics to evaluate bias in models trained on three datasets, finding that mitigation strategies can improve model performance for the unprivileged classes.

While numerous work have focused on define and mitigate bias, we have identified a gap in works that assess the direct impact of pre-training bias on the performance of machine learning algorithms, specifically in the Health domain. In this work, we aim to fill this gap by analyzing the correlation between pre-training bias metrics and the performance of supervised ML algorithms.

3. Methodology

Our methodology involves the following steps:

1. Data Collection: We collected four datasets from different sources
2. Data Preprocessing: We preprocessed the datasets to remove missing values, and encode categorical variables
3. Increase/Decrease Bias: We artificially increased and decreased bias in the datasets, creating a highly unbalanced and an equally balanced versions
4. Pre-training Metrics: We calculated the CI, KL Divergence, KS, and CDDL metrics for each dataset
5. Train Models: We trained supervised ML algorithms on the datasets and evaluated their performance using accuracy, precision, recall, and F1-score
6. Correlation Analysis: We analyzed the correlation between the pre-training bias metrics and the performance of the ML algorithms

The datasets, along with their source and protected attributes, are shown in Table 1.

Source	Dataset Name	Protected Attributes
[Newman et al. 1998]	HeartDataset	Sex
[Maslej et al. 2022]	IntersectionalBiasDataset	Sex, Race
[Teboul 2023]	DiabetesDataset	Sex, Age, Education, Income
[Tasci et al. 2022]	GliomaDataset	Gender, Race

Table 1. Datasets used in this work

4. Experiments and Results

5. Conclusion

6. References

References

- Bellamy, R. K. E., Dey, K., Hind, M., Hoffman, S. C., Houde, S., Kannan, K., Lohia, P., Martino, J., Mehta, S., Mojsilovic, A., Nagar, S., Ramamurthy, K. N., Richards, J., Saha, D., Sattigeri, P., Singh, M., Varshney, K. R., and Zhang, Y. (2018). Ai fairness 360: An extensible toolkit for detecting, understanding, and mitigating unwanted algorithmic bias. *arXiv*.
- Chen, I. Y., Pierson, E., Rose, S., Joshi, S., Ferryman, K., and Ghassemi, M. (2021). Ethical machine learning in healthcare. *Annual Review of Biomedical Data Science Annu. Rev. Biomed. Data Sci*, 2021:123–144.
- Hardt, M., Chen, X., Cheng, X., Donini, M., Gelman, J., Gollaprolu, S., He, J., Larroy, P., Liu, X., McCarthy, N., Rathi, A., Rees, S., Siva, A., Tsai, E., Vasist, K., Yilmaz, P., Zafar, M. B., Das, S., Haas, K., Hill, T., and Kenthapadi, K. (2021). Amazon sagemaker clarify: Machine learning bias detection and explainability in the cloud. *ArXiv*.
- Juhn, Y. J., Ryu, E., Wi, C.-I., King, K. S., Malik, M., Romero-Brufau, S., Weng, C., Sohn, S., Sharp, R. R., and Halamka, J. D. (2022). Assessing socioeconomic bias in machine learning algorithms in health care: a case study of the houses index. *Journal of the American Medical Informatics Association*, 29(7):1142–1151.
- Júnior, R. L. I., Silveira, L., de Faria, V. C. N., and Lorena, A. C. (2022). Justiça nas previsões de modelos de aprendizado de máquina: um estudo de caso com dados de reincidência criminal.
- Mandhala, V. N., Bhattacharyya, D., Midhunchakkaravarthy, D., and Kim, H. J. (2022). Detecting and mitigating bias in data using machine learning with pre-training metrics. *Ingenierie des Systemes d’Information*, 27:119–125.
- Maslej, M., Sikstrom, L., Reslan, D., and Wang, Y. (2022). Intersectional-bias-assessment.
- Newman, D., Hettich, S., Blake, C., and Merz, C. (1998). Uci repository of machine learning databases.
- Noseworthy, P. A., Attia, Z. I., Brewer, L. P. C., Hayes, S. N., Yao, X., Kapa, S., Friedman, P. A., and Lopez-Jimenez, F. (2020). Assessing and mitigating bias in medical

- artificial intelligence: The effects of race and ethnicity on a deep learning model for ecg analysis. *Circulation: Arrhythmia and Electrophysiology*.
- Obermeyer, Z., Powers, B., Vogeli, C., and Mullainathan, S. (2019). Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464):447–453.
- Park, Y., Hu, J., Singh, M., Sylla, I., Dankwa-Mullan, I., Koski, E., and Das, A. K. (2021). Comparison of methods to reduce bias from clinical prediction models of postpartum depression. *JAMA Network Open*, 4.
- Synthesized.io (2023). Fairlens: A toolkit for fair and interpretable machine learning. <https://github.com/synthesized-io/fairlens>. Accessed: 2024-10-12.
- Tasci, E., Zhuge, Y., Kaur, H., Camphausen, K., and Krauze, A. V. (2022). Hierarchical voting-based feature selection and ensemble learning model scheme for glioma grading with clinical and molecular characteristics. *International Journal of Molecular Sciences*, 23(22).
- Teboul, A. (2023). Diabetes health indicators dataset. Accessed: 2024-12-30.
- Zehlike, M., Castillo, C., Bonchi, F., Hajian, S., and Megahed, M. (2017). Fairness measures: Datasets and software for detecting algorithmic discrimination.
- Zhang, H., Lu, A. X., Abdalla, M., McDermott, M., and Ghassemi, M. (2020). Hurtful words. pages 110–120. Association for Computing Machinery, Inc.